

# Counting at Scale

APAM E4990

Modeling Social Data

Jake Hofman

Columbia University

February 10, 2017

# Previously

Claim:

Solving the **counting** problem **at scale** enables you to **investigate**  
many interesting **questions** in the **social sciences**

# Learning to count

Last week:

Counting at **small/medium** scales on a **single machine**

# Learning to count

Last week:

Counting at **small/medium** scales on a **single machine**

This week:

Counting at **large** scales **in parallel**

# What?



# What?

*“... to create building blocks for programmers who just happen to have **lots of data to store, lots of data to analyze, or lots of machines to coordinate**, and who don't have the time, the skill, or the inclination to become distributed systems experts to build the infrastructure to handle it.”*

-Tom White  
*Hadoop: The Definitive Guide*

# What?

Hadoop contains many subprojects:

- [Hadoop Common](#): The common utilities that support the other Hadoop subprojects.
- [Chukwa](#): A data collection system for managing large distributed systems.
- [HBase](#): A scalable, distributed database that supports structured data storage for large tables.
- [HDFS](#): A distributed file system that provides high throughput access to application data.
- [Hive](#): A data warehouse infrastructure that provides data summarization and ad hoc querying.
- [MapReduce](#): A software framework for distributed processing of large data sets on compute clusters.
- [Pig](#): A high-level data-flow language and execution framework for parallel computation.
- [ZooKeeper](#): A high-performance coordination service for distributed applications.

We'll focus on distributed computation with **MapReduce**.

# Who/when?

An overly brief history



# Who/when?

pre-2004

Doug Cutting and Mike Cafarella develop open source projects for web-scale indexing, crawling, and search



2004

## Dean and Ghemawat publish MapReduce programming model, used internally at Google

### MapReduce: Simplified Data Processing on Large Clusters

Jeffrey Dean and Sanjay Ghemawat

jeff@google.com, sanjay@google.com

*Google, Inc.*

#### Abstract


MapReduce is a programming model and an associated implementation for processing and generating large data sets. Users specify a *map* function that processes a key/value pair to generate a set of intermediate key/value pairs, and a *reduce* function that merges all intermediate values associated with the same intermediate key. Many real world tasks are expressible in this model, as shown in the paper.

given day, etc. Most such computations are conceptually straightforward. However, the input data is usually large and the computations have to be distributed across hundreds or thousands of machines in order to finish in a reasonable amount of time. The issues of how to parallelize the computation, distribute the data, and handle failures conspire to obscure the original simple computation with large amounts of complex code to deal with these issues.


# Who/when?

2006

Hadoop becomes official Apache project, Cutting joins Yahoo!,  
Yahoo adopts Hadoop



# Who/when?



## Where?



JPMORGAN CHASE &amp; CO.



**LinkedIn**

The New York Times



<http://wiki.apache.org/hadoop/PoweredBy>

# Why?

Why *yet another* solution?

(I already use too many languages/environments)

# Why?

Why a *distributed* solution?

(My desktop has TBs of storage and GBs of memory)

# Why?

Roughly how long to read 1TB from a commodity hard disk?



# Why?

Roughly how long to read 1TB from a commodity hard disk?

$$\frac{1 \text{ Gb}}{2 \text{ sec}} \times \frac{1 \text{ B}}{8 \text{ b}} \times 3600 \frac{\text{sec}}{\text{hr}} \approx 225 \frac{\text{GB}}{\text{hr}}$$

# Why?

Roughly how long to read 1TB from a commodity hard disk?<sup>1</sup>

$\approx 4\text{hrs}$

---

<sup>1</sup>SSDs provide a  $\sim 10\times$  speedup

# Why?

MAY 11, 2009

## Hadoop Sorts a Petabyte in 16.25 Hours and a Terabyte in 62 Seconds

We used **Apache Hadoop** to compete in **Jim Gray's Sort** benchmark. Jim's Gray's sort benchmark consists of a set of many related benchmarks, each with their own rules. All of the sort benchmarks measure the time to sort different numbers of 100 byte records. The first 10 bytes of each record is the key and the rest is the value. The **minute sort** must finish end to end in less than a minute. The **Gray sort** must sort more than 100 terabytes and must run for at least an hour. The best times we observed were:

Bytes	Nodes	Maps	Reduces	Replication	Time
500,000,000,000	1406	8000	2600	1	59 seconds
1,000,000,000,000	1460	8000	2700	1	62 seconds
100,000,000,000,000	3452	190,000	10,000	2	173 minutes
1,000,000,000,000,000	3658	80,000	20,000	2	975 minutes

<http://bit.ly/petabytesort>

# Typical scenario

## Store, parse, and analyze high-volume server logs,

```
[16/May/2010:07:28:49 -0400] "GET /autonomous_css/style.css HTTP/1.1" 200 2806 "http://www.jakehofman.com/" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:49 -0400] "GET /cleanlooks/test3.gif HTTP/1.1" 404 339 "http://www.jakehofman.com/" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:49 -0400] "GET / HTTP/1.1" 200 16458 "http://www.technologyreview.com/communications/25326/?a=f&utm_source=feedburner&utm_medium=feed&utm_campaign=Feed%3A+YDNLlinkBlog%28Yahoo%21+Developer+Network+Linkblog%29" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:50 -0400] "GET /favicon.ico HTTP/1.1" 404 330 "-" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:57 -0400] "GET /autonomous_css/style.css HTTP/1.1" 304 - "http://www.jakehofman.com/" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:57 -0400] "GET /cleanlooks/test3.gif HTTP/1.1" 404 339 "http://www.jakehofman.com/" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:56 -0400] "GET / HTTP/1.1" 200 16458 "http://www.technologyreview.com/communications/25326/?a=f&utm_source=feedburner&utm_medium=feed&utm_campaign=Feed%3A+YDNLlinkBlog%28Yahoo%21+Developer+Network+Linkblog%29" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
```

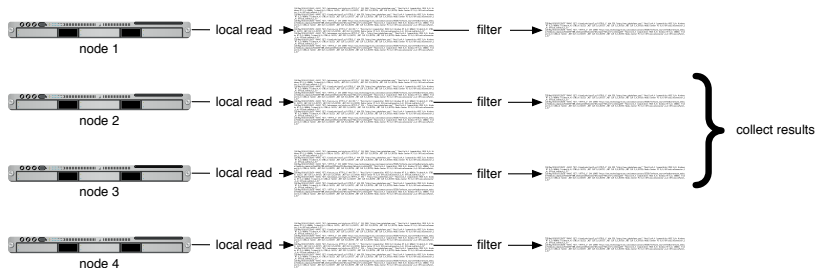
e.g. how many search queries match “icwsm”?

# MapReduce: 30k ft

Break large problem into smaller parts, solve in parallel, combine results

# Typical scenario

“Embarassingly parallel”  
(or nearly so)

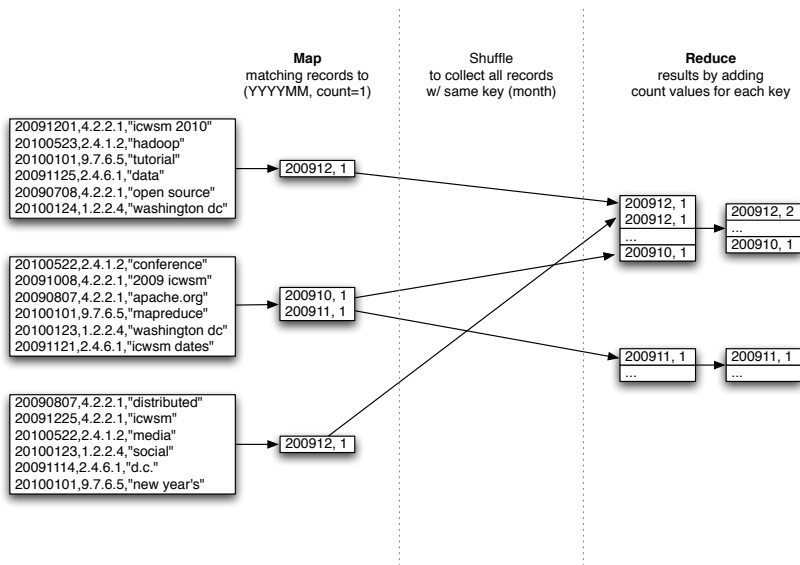


# Typical scenario++

## How many search queries match “icwsm”, grouped by month?

```
[16/May/2010:07:28:49 -0400] "GET /autonomous_css/style.css HTTP/1.1" 200 2806 "http://www.jakehofman.com/" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:49 -0400] "GET /cleanlooks/test3.gif HTTP/1.1" 404 339 "http://www.jakehofman.com/" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:49 -0400] "GET / HTTP/1.1" 200 16458 "http://www.technologyreview.com/communications/25326/?a=feed&utm_source=Feedburner&utm_medium=feed&utm_campaign=Feed%3A+YDNLlinkBlog%3A28Yahoo%3A221+Developer+Network+linkblog%3A29" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:50 -0400] "GET /favicon.ico HTTP/1.1" 404 330 "-" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:57 -0400] "GET /autonomous_css/style.css HTTP/1.1" 304 - "http://www.jakehofman.com/" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:57 -0400] "GET /cleanlooks/test3.gif HTTP/1.1" 404 339 "http://www.jakehofman.com/" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
[16/May/2010:07:28:56 -0400] "GET / HTTP/1.1" 200 16458 "http://www.technologyreview.com/communications/25326/?a=feed&utm_source=Feedburner&utm_medium=feed&utm_campaign=Feed%3A+YDNLlinkBlog%3A28Yahoo%3A221+Developer+Network+linkblog%3A29" "Mozilla/4.0 (compatible; MSIE 8.0; Windows NT 6.1; WOW64; Trident/4.0; GTB6.4; SLCC2; .NET CLR 2.0.50727; .NET CLR 3.5.30729; .NET CLR 3.0.30729; Media Center PC 6.0; OfficeLiveConnector.1.4; OfficeLivePatch.1.3)"
```

# MapReduce: example





# MapReduce: paradigm

Programmer specifies **map** and **reduce** functions

# MapReduce: paradigm

**Map:** transforms input record to intermediate (key, value) pair

```
def mapper(record):  
    # input: a single record  
  
    # parse / transform / filter record  
    ...  
  
    # output: intermediate key(s) and value(s)  
    output( (key, value) )
```

# MapReduce: paradigm

Shuffle: collects all intermediate records by key

Record assigned to reducers by  $\text{hash}(\text{key}) \% \text{num\_reducers}$

Reducers perform a **merge sort** to collect records with same key

# MapReduce: paradigm

Reduce: transforms all records for given key to final output

```
def reducer(key, records):  
    # input: intermediate key and all values  
  
    # initialize variables, e.g. counters  
  
    for record in records:  
        # parse record  
        ...  
  
        # update variables, e.g. count  
  
    # output: final key(s) and value(s)  
    output( (key, value) )
```

# MapReduce: paradigm

Distributed read, shuffle, and write are transparent to programmer

# MapReduce: principles

- Move **code to data** (local computation)
- Allow programs to **scale transparently** w.r.t size of input
- **Abstract away** fault tolerance, synchronization, etc.

# MapReduce: strengths

- Batch, offline jobs
- Write-once, read-many across full data set
- Usually, though not always, simple computations
- I/O bound by disk/network bandwidth

# !MapReduce

What it's not:

- High-performance parallel computing, e.g. MPI
- Low-latency random access relational database
- Always the right solution



# Word count

the quick brown fox  
jumps over the lazy dog  
who jumped over that  
lazy dog -- the fox ?



dog	2
--	1
the	3
brown	1
fox	2
jumped	1
lazy	2
jumps	1
over	2
quick	1
that	1
who	1
?	1

# Word count

**Map:** for each line, output each word and count (of 1)

```
the quick brown fox
-----
jumps over the lazy dog
-----
who jumped over that
-----
lazy dog -- the fox ?
```



```
the 1
quick 1
brown 1
fox 1
-----
jumps 1
over 1
the 1
lazy 1
dog 1
-----
who 1
jumped 1
over 1
-----
that 1
lazy 1
dog 1
-- 1
the 1
fox 1
? 1
```

# Word count

Shuffle: collect all records for each word

the quick brown fox  
-----  
jumps over the lazy dog  
-----  
who jumped over that  
-----  
lazy dog -- the fox ?




-- 1  
-----  
? 1  
-----  
brown 1  
-----  
dog 1  
dog 1  
-----  
fox 1  
fox 1  
-----  
jumped 1  
-----  
jumps 1  
-----  
lazy 1  
lazy 1  
-----  
over 1  
over 1  
-----  
quick 1  
-----  
that 1  
-----  
the 1  
the 1  
the 1  
-----  
who 1

# Word count

**Reduce:** add counts for each word

--	1	
-----		
?	1	
-----		
brown	1	
-----		
dog	1	
dog	1	
-----		
fox	1	
fox	1	
-----		
jumped	1	
-----		
jumps	1	
-----		
lazy	1	
lazy	1	
-----		
over	1	
over	1	
-----		
quick	1	
-----		
that	1	
-----		
the	1	
the	1	
the	1	
-----		
who	1	



--	1	
?	1	
brown	1	
dog	2	
fox	2	
jumped	1	
jumps	1	
lazy	2	
over	2	
quick	1	
that	1	
the	3	
who	1	

# Word count

the quick brown fox  
-----  
jumps over the lazy dog  
-----  
who jumped over that  
-----  
lazy dog -- the fox ?



dog 1  
dog 1  
-----  
-- 1  
-----  
the 1  
the 1  
the 1  
-----  
brown 1  
-----  
fox 1  
fox 1  
-----  
jumped 1  
-----  
lazy 1  
lazy 1  
-----  
jumps 1  
-----  
over 1  
over 1  
-----  
quick 1  
-----  
that 1  
-----  
? 1  
-----  
who 1



dog 2  
-- 1  
the 3  
brown 1  
fox 2  
jumped 1  
lazy 2  
jumps 1  
over 2  
quick 1  
that 1  
who 1  
? 1

# WordCount.java

```
1. package org.myoung;
2.
3. import java.io.IOException;
4. import java.util.*;
5.
6. import org.apache.hadoop.fs.Path;
7. import org.apache.hadoop.conf.*;
8. import org.apache.hadoop.io.*;
9. import org.apache.hadoop.mapred.*;
10. import org.apache.hadoop.util.*;
11.
12. public class WordCount {
13.
14.     public static class Map extends MapReduceBase implements Mapper<LongWritable, Text, Text, IntWritable> {
15.         private final static IntWritable one = new IntWritable(1);
16.         private Text word = new Text();
17.
18.         public void map(LongWritable key, Text value, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
19.             String line = value.toString();
20.             StringTokenizer tokenizer = new StringTokenizer(line);
21.             while (tokenizer.hasMoreTokens()) {
22.                 word.set(tokenizer.nextToken());
23.                 output.collect(word, one);
24.             }
25.         }
26.     }
27.
28.     public static class Reduce extends MapReduceBase implements Reducer<Text, IntWritable, Text, IntWritable> {
29.         public void reduce(Text key, Iterator<IntWritable> values, OutputCollector<Text, IntWritable> output, Reporter reporter) throws IOException {
30.             int sum = 0;
31.             while (values.hasNext()) {
32.                 sum += values.next().get();
33.             }
34.             output.collect(key, new IntWritable(sum));
35.         }
36.     }
37.
38.     public static void main(String[] args) throws Exception {
39.         JobConf conf = new JobConf(WordCount.class);
40.         conf.setJobName("wordcount");
41.
42.         conf.setOutputKeyClass(Text.class);
43.         conf.setOutputValueClass(IntWritable.class);
44.
45.         conf.setMapperClass(Map.class);
46.         conf.setCombinerClass(Reduce.class);
47.         conf.setReducerClass(Reduce.class);
48.
49.         conf.setInputFormat(TextInputFormat.class);
50.         conf.setOutputFormat(TextOutputFormat.class);
51.
52.         FileInputFormat.setInputPaths(conf, new Path(args[0]));
53.         FileOutputFormat.setOutputPath(conf, new Path(args[1]));
54.
55.         JobClient.runJob(conf);
56.     }
57. }
58. }
```

# Hadoop streaming

Hadoop streaming is a utility that comes with the Hadoop distribution. The utility allows you to create and run map/reduce jobs with any executable or script as the mapper and/or the reducer. For example:

```
$HADOOP_HOME/bin/hadoop jar $HADOOP_HOME/hadoop-streaming.jar \  
-input myInputDirs \  
-output myOutputDir \  
-mapper /bin/cat \  
-reducer /bin/wc
```

# Hadoop streaming

MapReduce for \*nix geeks<sup>2</sup>:

```
# cat data | map | sort | reduce
```

Where the sort is a hack to approximate a distributed group-by:

- Mapper reads input data from stdin
- Mapper writes output to stdout
- Reducer receives input, grouped by key, on stdin
- Reducer writes output to stdout

---

<sup>2</sup><http://bit.ly/michaelnoll>



Locally:

```
# cat data | tr " " "\n" | sort | uniq -c
```

# wordcount.sh

Locally:

```
# cat data | tr " " "\n" | sort | uniq -c
```



Distributed:

```
$HADOOP_HOME/bin/hadoop jar $HADOOP_HOME/hadoop-streaming.jar \  
-input README.txt \  
-output wordcount \  
-mapper 'tr " " "\n" \  
-reducer 'uniq -c'
```

# Transparent scaling

Use **the same code** on **MBs locally** or **TBs across thousands of machines**.

# wordcount.py

```
from hstream import HStream
import sys
import re
from collections import defaultdict

class WordCount(HStream):

    def mapper(self, record):
        for word in " ".join(record).split():
            self.write_output((word, 1))

    def reducer(self, key, records):
        total = 0

        for record in records:
            word, count = record

            total += int(count)

        self.write_output( (word, total) )

if __name__ == '__main__':
    WordCount()
```